Information Extraction: A short introduction

Rafael Carrascosa

Machinalis

rcarrascosa@machinalis.com

November 4, 2014

About the speaker

- Rafael Carrascosa.
- Background in Computer Science.
- Works as developer and Machine Leaning specialist for Machinalis.
- Machinalis (www.machinalis.com) is a software research and development company based on Argentina.
- Leaded the delopment of IEPY.

Overview

1. Information Extraction
   - What is IE?
   - Applications of IE
   - Limitations of IE
   - How IE is done, a sketch
   - Reasoning and Knowledge Bases
2. IE with IEPY
   - Overview
   - Relation extraction with IEPY
   - Corpora construction

What is Information Extraction

In one line

Is about getting information from text automatically

- Main activity: extract **relation** instances between **entities**.
- **Entity**: Any-thing: A person, a place, an organization, a date, etc.
- **Relation**: Can be any kind of linking concept between entities.
- Also posible but not common: IE over audio, images or video.

An example of entities and relations

Example text

Harvard University is a private Ivy League research university in **Cambridge**, Massachusetts, established 1636.

- `Harvard University`: Is an **entity** of kind organization.
- `Cambridge`: Is an **entity** of kind location.
- *Located*(X, Y) is a **relation** that links organizations with places.

An example of entities and relations

Example text

Harvard University is a private Ivy League research university in **Cambridge**, Massachusetts, established 1636.

- `Harvard University`: Is an **entity** of kind organization.
- `Cambridge`: Is an **entity** of kind location.
- *Located*(X, Y) is a **relation** that links organizations with places.

IE deals with detecting that

$$\textit{Located}(\text{Harvard University}, \text{Cambridge})$$

is a **relation instance** in the text.

Example application: Protests

Example application: Protests

Use news text to build a map of events like politic acts or public protests.

Example application: Protests

Use news text to build a map of events like politic acts or public protests.

- News sites, social networks
- $happened\_at(protest, date)$
- $happened\_in(protest, location)$
- Data can end up in a map, or a timeline.

Example application: Foodborne outbreaks

Use social media comments to detect foodborne outbreaks.

Example application: Foodborne outbreaks

Use social media comments to detect foodborne outbreaks.

- $ate\_at(person, restaurant)$
- $has\_symptom(person, symptom)$
- True story: a colaboration between Yelp, Columbia University and NYC's Department of Health and Mental Hygiene.

More example applications

- Starting from business news build a network graph with deals, investments, acquisitions.
- With sanitary documents, build a timeline of diseases, symptoms and treatments.
- Reconstructing state terrorism victims' fate from military documents.
- Use social media to measure the growth of statup incubators and co-working spaces.
- ...

What is not IE

What is not IE

Is not relation discovery, relations have to be known a-priori.

What is not IE

Is not relation discovery, relations have to be known a-priori.

Is not open domain, IE is domain specific.

A sketch of an IE system

Steps

- Basic natural language processing
- Entity recognition and linking
- Relation extraction

A sketch of an IE system

Steps

- Basic natural language processing
- Entity recognition and linking
  - Named entity recognition
  - Coreference resolution
  - Name linking
- Relation extraction

A sketch of an IE system

Steps

- Basic natural language processing
- Entity recognition and linking
  - Named entity recognition
  - Coreference resolution
  - Name linking
- Relation extraction
  - Rule-based
  - Corpus-based

Relation extraction

Rule-based

A regular expression to detect the presence of a relation.

- Handcrafted by an expert
- Suggested and validated by an algorithm

Corpus-based

For each pair of entities, a binary classification:

- Classifiers: Naive bayes, support vector machines, etc.
- Language models: Markov models, conditional random fields, etc.

Knowledge Bases and Reasoning

Knowledge bases

- Freebase
- Linked open data

Knowledge Bases and Reasoning

Knowledge bases

- Freebase
- Linked open data

Reasoning

Semantic reasoners, they exist.

IE summary

- Getting information out of large amounts of text documents.
- Extract **relation** instances between **entities**.
- Has many practical applications.
- Domain-specific.
- Solved using rules or statistic tools.

IE with IEPY

What is IEPY

- Tool for Information Extraction.
- Open source (BSD).
- Written in Python.
- Developed by Machinalis.
- Aimed at:
  - IE users
  - IE scientists
- `https://github.com/machinalis/iepy`

Features

For IE users

- NER and coreference resolution.
- **Active learning** relation extraction.
- **Rule-based relation extraction**.

For IE scientists

- **Corpus annotation** tool.
- Web-based UI.
- Easily hackable.

IEPY: Active learning core

Statistical classifier

- Support Vector Machine (by default).
- Features: Bag-of-words, POS tags, entity distance, etc.
- Active learning.
- Tunable to high-precision or high-recall.
- Web-based UI to interact with the expert.
- Easily hackable (`scikit-klearn` and `django`)

IEPY: Active learning core

Active learning goal

Minimize human effort

IEPY: Active learning core

Active learning goal

Minimize human effort

Idea

- Query a human expert.
- Use the answer to compute the next most useful questions to ask.
- Repeat.

Uses less human time to achieve the same performance.

IEPY: Active learning core

Lee entered the **Slade School of Art** in **1911** where **he** became friendly with **Robert Gibbings** and **Paul Nash** .

Complete:

Lee was born **1911** ? — Skipped labeling of this evidence

he was born **1911** ? — Skipped labeling of this evidence

Robert Gibbings was born **1911** ? — Skipped labeling of this evidence

Paul Nash was born **1911** ? — Skipped labeling of this evidence

IEPY: Rule core

Rule-based

- Enhanced regular expressions.
- Expresiveness: Words, POS tags, entities, entity kinds, etc.
- Positive or negative rules.
- Have to be written in Python.
- Simplified syntax

Rule-based

```
Subject + Token(", born") + Object + anything
```

Rule-based

```
Subject + Token(", born") + Object + anything
```

Matches something like

Lyle Eugene Hollister, born **6 July 1923** in Sioux Falls...

Rule-based

```
anything + Subject + Token("was born") + Pos("IN") + Object + anything
```

Rule-based

```
anything + Subject + Token("was born") + Pos("IN") + Object + anything
```

Matches something like

. . . **Shamsher M. Chowdhury** was born in **1950** . . .

Corpora construction

Evaluation?

Eventually you'll need to evaluate your IE performance.

Corpora construction

Evaluation?

Eventually you'll need to evaluate your IE performance.

- How much noise has the information you extracted?
- How many relevant facts you are leaving behind?

Corpora construction

IEPY has a corpora construction tool

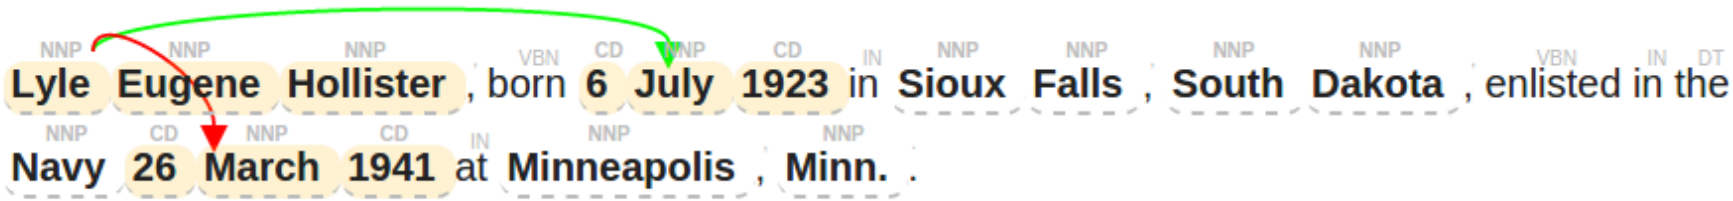

Since is a web-based it allows multiple experts to work simultaneously

Some performance figures

Active learning

On an easy relation (date of birth)

86% prec. and 80% rec.

with an hour's worth of labeling.

Rule based

On an easy relation (date of birth)

98.65% prec. and 38% rec.

investing 4 hours (11 rules).

Corpora construction

Rate of 50 documents per hour (we labeled 4300 docs).

The end

That's it, thank you!